

Code-along 03

FirstName LastName

Setup

Packages

Load the standard packages.

```
library(here)
library(tidyverse)
library(haven) # not core tidyverse
library(gssr)
library(gssrdoc)
library(summarytools)
```

Load your data & codebook

```
# Get the data all survey years
data(gss_all)

# Load the codebook
data(gss_dict)
```

Data Management I

The pipe |>

The pipe operator passes what comes before it into the function that comes after it as the first argument in that function.

```
sum(1, 2)
```

```
[1] 3
```

```
1 |>  
  sum(2)
```

```
[1] 3
```

dplyr grammar

What's the advantage of **dplyr** grammar? We can sequence data manipulation!

```
gss_all |>  
  select(year, sex, agekdbnr) |>  
  filter(year == 2022) |>  
  drop_na(sex, agekdbnr) |>  
  group_by(sex) |>  
  summarise(avg = mean(agekdbnr))
```

```
# A tibble: 2 x 2  
  sex      avg  
  <dbl+lbl> <dbl>  
1 1 [male]  26.3  
2 2 [female] 23.7
```

select(), filter(), and drop_na()

Use **select()** to pick specific columns from your dataset.

Use **filter()** to keep rows that meet a condition.

Use **drop_na()** to remove rows with missing (NA) values.

```
gss_all |>  
  select(year, sex, agekdbnr) |>  
  filter(year == 2022) |>  
  drop_na(sex, agekdbnr) |>
```

```
Error in parse(text = input): <text>:6:0: unexpected end of input
```

```
4:   drop_na(sex, agekdbnr) |>
```

```
5:
```

```
~
```

group_by() and summarize()

Use `group_by()` to organize your data into groups based on one or more variables.

Use `summarize()` to compute statistics like total, mean, or median for each group.

```
gss_all |>
  select(year, sex, agekdbnr) |>
  filter(year == 2022) |>
  drop_na(sex, agekdbnr) |>
  group_by(sex) |>
  summarise(freq = n())
```

①
②
③
④
⑤
⑥

- ① Start with the `gss_all` data frame:
- ② Keep only the variables in the dataset that we need.
- ③ Keep only the respondents from the 2022 survey
- ④ Remove any observations with missing data for our key variables
- ⑤ Do the next steps separately for each group in the variable
- ⑥ Creates a new data frame with one row for each combination of grouping variables

```
# A tibble: 2 x 2
  sex      freq
  <dbl> <int>
1 1 [male]  1031
2 2 [female] 1363
```

dplyr() in action

Compare the average and median age at first childbirth for U.S. men and women in 2022.

```
1 gss_all |>
2   select(year, sex, agekdbnr) |>
3   filter(year == 2022) |>
4   drop_na(sex, agekdbnr) |>
5   group_by(sex) |>
6   summarise(
7     freq = n(),
8     avg = mean(agekdbnr),
9     med = median(agekdbnr)
10  )
```

⑤

```
# A tibble: 2 x 4
  sex      freq  avg  med
<dbl>+<lbl> <int> <dbl> <dbl>
1 1 [male]   1031  26.3   25
2 2 [female] 1363  23.7   23
```

mutate() in action

Use `mutate()` to add new columns or change existing ones.

What proportion of new parents were teenagers (e.g., under 18 years old)?

```
1 gss_all |>
2   select(year, agekdbnr) |>
3   filter(year == 2022) |>
4   drop_na(agekdbnr) |>
5   mutate(teen_parent = (agekdbnr < 18) * 1) |>
6   summarise(proportion = mean(teen_parent))
```

- ① Manipulate the dataset to variables and observations of interest
- ② Remove missing or `summarise()` will report NA
- ③ Returns a logical vector (TRUE, FALSE, or NA) and `* 1` converts it to numeric
- ④ Mean of the variable, which equals the proportion of 1s

```
# A tibble: 1 x 1
  proportion
    <dbl>
1    0.0773
```

Use `case_when()` inside `mutate()` to create values based on conditions.

What proportion of new parents had their first child as teenagers, in their 20s, 30s, or after age 40?

```
1 gss_all <- gss_all |>
2   mutate(age_groups = case_when(
3     agekdbnr < 18 ~ "<18",
4     agekdbnr >= 18 & agekdbnr <= 29 ~ "18-29",
5     agekdbnr >= 30 & agekdbnr <= 39 ~ "30-39",
6     agekdbnr >= 40 ~ "40+",
7     TRUE ~ NA_character_)
8 )
```

```

9
10 gss_all |>
11   filter(year == 2022) |>
12   freq(age_groups, report.nas = FALSE, headings = FALSE)

```

	Freq	%	% Cum.
<18	186	7.73	7.73
18-29	1704	70.82	78.55
30-39	463	19.24	97.80
40+	53	2.20	100.00
Total	2406	100.00	100.00

Assignment operators

Let's make a tiny data frame to use as an example:

```

df <- tibble(x = c(1, 2, 3, 4, 5), y = c("a", "a", "b", "c", "c"))
df

```

```

# A tibble: 5 x 2
      x y
  <dbl> <chr>
1     1 a
2     2 a
3     3 b
4     4 c
5     5 c

```

```

df |>
  mutate(x = x * 2)

```

```

# A tibble: 5 x 2
      x y
  <dbl> <chr>
1     2 a
2     4 a
3     6 b
4     8 c
5    10 c

```

```
df
```

```
# A tibble: 5 x 2
      x y
  <dbl> <chr>
1     1 a
2     2 a
3     3 b
4     4 c
5     5 c
```

```
#| label: assignment

df <- df |>
  mutate(x = x * 2)
```

```
df
```

```
# A tibble: 5 x 2
      x y
  <dbl> <chr>
1     2 a
2     4 a
3     6 b
4     8 c
5    10 c
```

Do something, save result, overwriting original

```
df <- tibble(
  x = c(1, 2, 3, 4, 5),
  y = c("a", "a", "b", "c", "c")
)

df <- df |>
  mutate(x = x * 2)

df
```

```
# A tibble: 5 x 2
      x y
  <dbl> <chr>
1     2 a
2     4 a
3     6 b
4     8 c
5    10 c
```

Do something, save result, *not* overwriting original

```
df <- tibble(
  x = c(1, 2, 3, 4, 5),
  y = c("a", "a", "b", "c", "c")
)

df_new <- df |>
  mutate(x = x * 2)

df_new
```

```
# A tibble: 5 x 2
      x y
  <dbl> <chr>
1     2 a
2     4 a
3     6 b
4     8 c
5    10 c
```

Do something, save result, overwriting original when you shouldn't

```
df <- tibble(
  x = c(1, 2, 3, 4, 5),
  y = c("a", "a", "b", "c", "c")
)

df <- df |>
  group_by(y) |>
  summarize(mean_x = mean(x))

df
```

```
# A tibble: 3 x 2
  y      mean_x
<chr> <dbl>
1 a         1.5
2 b          3
3 c         4.5
```

Do something, save result, not overwriting original when you shouldn't

```
df <- tibble(
  x = c(1, 2, 3, 4, 5),
  y = c("a", "a", "b", "c", "c")
)

df_summary <- df |>
  group_by(y) |>
  summarize(mean_x = mean(x))

df_summary
```

```
# A tibble: 3 x 2
  y      mean_x
<chr> <dbl>
1 a         1.5
2 b          3
3 c         4.5
```

Do something, save result, overwriting original data frame

```
df <- tibble(
  x = c(1, 2, 3, 4, 5),
  y = c("a", "a", "b", "c", "c")
)

df <- df |>
  mutate(z = x + 2)

df
```

```
# A tibble: 5 x 3
  x y      z
<dbl> <chr> <dbl>
1 1 a      3
```


2	2 a	4
3	3 b	5
4	4 c	6
5	5 c	7

Do something, save result, overwriting original column

```
df <- tibble(
  x = c(1, 2, 3, 4, 5),
  y = c("a", "a", "b", "c", "c")
)
df <- df |>
  mutate(x = x + 2)
df
```

A tibble: 5 x 2

	x	y
	<dbl>	<chr>
1	3	a
2	4	a
3	5	b
4	6	c
5	7	c

Do something, save result, not overwriting original.

```
1 gss_all <- gss_all |>
2   mutate(age_groups = case_when(
3     agekdbrn < 18 ~ "<18",
4     agekdbrn >= 18 & agekdbrn <= 29 ~ "18-29",
5     agekdbrn >= 30 & agekdbrn <= 39 ~ "30-39",
6     agekdbrn >= 40 ~ "40+",
7     TRUE ~ NA_character_)
8   )
9
10 gss_all |>
11   filter(year == 2022) |>
12   freq(age_groups, report.nas = FALSE, headings = FALSE)
```

Do something and show me

```

1 gss_all |>
2   select(year, agekdbnr) |>
3   filter(year == 2022) |>
4   drop_na(agekdbnr) |>
5   mutate(age_groups = case_when(
6     agekdbnr < 18 ~ "<18",
7     agekdbnr >= 18 & agekdbnr <= 29 ~ "18-29",
8     agekdbnr >= 30 & agekdbnr <= 39 ~ "30-39",
9     agekdbnr >= 40 ~ "40+",
10    TRUE ~ NA_character_)) |>
11   group_by(age_groups) |>
12   summarise(
13     count = n(),
14     proportion = round(count / sum(count), 3)
15   )

```

Summary Statistics

Median & Mode

Let's use dplyr grammar to find the median and mode for the `childs` variable.

```

gss_all$childs <- zap_missing(gss_all$childs)
gss_all$childs <- as_factor(gss_all$childs)
gss_all$childs <- droplevels(gss_all$childs)

```

```

gss_all |>
  filter(year == 2024) |>
  freq(childs, report.nas = FALSE) |>
  tb()

```

①

②

③

- ① Use `dplyr` grammar, starting with the name of the df and a pipe
- ② Use the `freq()` function as usual
- ③ Add the `tb()` function to turn the table into a `tibble`

```

# A tibble: 9 x 4
  childs    freq    pct pct_cum
<fct>    <dbl> <dbl> <dbl>
1 0      1029 31.4    31.4
2 1       484 14.8    46.2

```

3	2	851	26.0	72.1
4	3	475	14.5	86.6
5	4	243	7.41	94.0
6	5	96	2.93	96.9
7	6	53	1.62	98.6
8	7	16	0.488	99.1
9	8 or more	31	0.946	100

Median & Mean

Let's use dplyr grammar to find the median and mean for the `hrs1` variable.

```
median(gss_all$hrs1, na.rm=TRUE)
mean(gss_all$hrs1, na.rm=TRUE)

# show me summary statistics
summary(gss_all$hrs1)
```

①

- ① `na.rm` is a logical evaluating to TRUE or FALSE indicating whether NA values should be stripped before the computation proceeds.

```
[1] 40
[1] 41.11279
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
  0.00   37.00   40.00   41.11   48.00   89.00  32371
```

descr()

Univariate statistics for numerical data

```
gss_all |>
  filter(year == 2024) |>
  drop_na(hrs1) |>
  descr(hrs1,
        stats = "common") |>
  tb()
```

①

②

- ① Which stats to produce. Either “all” (default), “fivenum”, “common” (see Details), or a selection. See `?descr`
- ② Don't forget to make a tidy dataset out of `freq()` or `descr()` outputs

```
# A tibble: 1 x 9
  variable mean    sd  min  med  max n.valid    n pct.valid
  <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>
1 year      39.4  13.9    0   40   89   1768  1768    100
```

```
summarize()
```

```
1 gss_all |>
2   select(year, hrs1, sex) |>
3   filter(year == 2024) |>
4   drop_na(hrs1, sex) |>
5   group_by(as_factor(sex)) |>
6   summarise(
7     count = n(),
8     min = min(hrs1),
9     median = median(hrs1),
10    max = max(hrs1),
11    mean = round(mean(hrs1), digits = 2),
12    sd = sd(hrs1)
13  )
```

```
# A tibble: 2 x 7
  `as_factor(sex)` count min      median max      mean    sd
  <fct>           <int> <dbl+lbl> <dbl> <dbl+lbl> <dbl> <dbl>
1 male             869 0          40 89 [89+ hours] 41.7 13.7
2 female           891 0          40 89 [89+ hours] 37.3 13.7
```

Think Like a Statistician

On average, in 2024, did parents with 4 or more kids work fewer hours for pay than other parents?

How do we find out?

- Make `childs` a numeric variable
- `filter()` the data to only 2024 respondents
- `select()` only the variables you need: `year`, `childs`, `hrs`
- use `mutate()` to create a character variable with 4 categories: 1 child, 2 children, 3 children, 4 or more children (and `TRUE ~ NA_character_` for missing data)
- use `drop_by()` to remove missing data for your new variable and `hrs1`

- use `group_by()` to group the data by your new variable
- use `summarise()` to create count, mean, median, and sd summary statistics

```
# TIP: It's often easier to play with your code in an R script first.  
# Then, copy and past your working R code into this code-chunk.
```

Your Data Take

What's your conclusion to our initial research question?