

Code-along 09

FirstName LastName

Setup

Packages

Load the standard packages and new package `modelbased`.

```
# install.packages("marginaleffects") # needed for modelbased
# install.packages("modelbased")

library(tidyverse)
library(haven) # not core tidyverse
library(gssr)
library(gssrdoc)
library(summarytools)
library(gtsummary)
library(moderndiver)
library(modelbased)
```

Load your data

```
# Load all gss
data(gss_all)
```

Make a df with only the (pretty) categorical and continuous variables we'll analyze.

```

# Categorical Variables
my_cat <- gss_all |>
  select(id, year, polviews, sex, marital) |>
  zap_missing() |>
  mutate(
    polviews = as_factor(polviews),
    sex = as_factor(sex),
    marital = as_factor(marital)) |>
  droplevels()

# Continuous Variables
my_con <- gss_all |>
  select(id, year, tvhours, age, educ, hrs1, realrinc, childs) |>
  mutate(
    age = as.numeric(age),
    educ = as.numeric(educ),
    hrs1 = as.numeric(hrs1),
    realrinc = as.numeric(realrinc),
    childs = as.numeric(childs))

# Combine the two dataframes
my_data <- full_join(my_cat, my_con, by = c("year", "id")) # match by 2 vars

```


Two continuous variables

Have the average number of children changed over time?

```

# Create scatterplot with a regression line
my_data |>
  group_by(year) |>
  summarize(avg = mean(childs, na.rm = TRUE)) |>
  ggplot(aes(x = year, y = avg)) +
  geom_point(color = "#18BC9C", size = 3) +
  geom_smooth(method = "lm", se = FALSE, color = "#3498DB") +
  labs(
    title = "Scatterplot of Average Number of Children",
    x = "Year",
    y = "Average number of children") +
  theme_minimal()

```



Pearson's Correlation (R)

```
cor.test(
  ~ childs + year,
  data = my_data,
  method = "pearson",
  na.action = na.omit)
```

Pearson's product-moment correlation

```
data:  childs and year
t = -18.896, df = 75405, p-value < 2.2e-16
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 -0.07575205 -0.06154441
sample estimates:
      cor
-0.06865171
```

Simple Regression

```
model_01 <- lm(formula = childs ~ year, data = my_data)
get_regression_table(model_01)
```

```
# A tibble: 2 x 7
  term      estimate std_error statistic p_value lower_ci upper_ci
  <chr>      <dbl>    <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 intercept  17.2      0.811     21.2     0     15.7    18.8
2 year      -0.008     0      -18.9     0    -0.008 -0.007
```

PRE Measure

Square R

```
my_test <- cor.test(
  ~ childs + year,
  data = my_data,
  method = "pearson",
  na.action = na.omit
)

my_test$estimate^2
```

```
cor
0.004713057
```


Get r^2

```
model_01 <- lm(formula = childs ~ year, data = my_data)
get_regression_summaries(model_01)
```

```
# A tibble: 1 x 9
  r_squared adj_r_squared  mse  rmse sigma statistic p_value  df  nob
  <dbl>      <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl> <dbl> <dbl>
1  0.005      0.005  3.07  1.75  1.75   357.     0    1 75407
```

Are the trends explained by increases in education?

```
# Create scatterplot with regression line
my_data |>
  group_by(year) |>
  summarize(avg = mean(educ, na.rm = TRUE)) |>
  ggplot(aes(x = year, y = avg)) +
  geom_point(color = "#18BC9C", size = 3) +
  geom_smooth(method = "lm", se = FALSE, color = "#3498DB") +
  labs(
    title = "Scatterplot of Average Years of Education",
    x = "Year",
    y = "Average years of education"
  ) +
  theme_minimal()
```



```
# Create scatterplot with regression line
my_data |>
  group_by(educ) |>
  summarize(avg = mean(childs, na.rm = TRUE)) |>
  ggplot(aes(x = educ, y = avg)) +
  geom_point(color = "#18BC9C", size = 3) +
  geom_smooth(method = "lm", se = FALSE, color = "#3498DB") +
  labs(
```


```

  title = "Scatterplot of Average Number of Children",
  x = "Average years of education",
  y = "Average number of children"
) +
theme_minimal()

```

Warning: Removed 1 row containing non-finite outside the scale range (`stat\_smooth()`).

Warning: Removed 1 row containing missing values or values outside the scale range (`geom\_point()`).



Multiple Regression

```

model_02 <- lm(formula = childs ~ year + educ, data = my_data)
get_regression_table(model_02)

```

```

# A tibble: 3 x 7
  term          estimate std_error statistic p_value lower_ci upper_ci
  <fct>          <dbl>    <dbl>    <dbl>   <dbl>   <dbl>
1 year          0.000    0.000     0.000   1.000    0.000  0.000
2 educ          0.000    0.000     0.000   1.000    0.000  0.000
3 (Intercept)  1.500    0.000    1500.0   0.000    1.500  1.500

```

Characteristic	Beta	95% CI	p-value
(Intercept)	5.3	3.7, 6.9	<0.001
GSS year for this respondent	0.00	0.00, 0.00	0.045
educ	-0.13	-0.13, -0.13	<0.001

Abbreviation: CI = Confidence Interval

	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	intercept	5.26	0.813	6.47	0	3.66	6.85
2	year	-0.001	0	-2.00	0.045	-0.002	0
3	educ	-0.131	0.002	-64.5	0	-0.134	-0.127

PRE Measure

```
get_regression_summaries(model_02)
```

A tibble: 1 x 9

	r_squared	adj_r_squared	mse	rmse	sigma	statistic	p_value	df	nobs
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.057	0.057	2.90	1.70	1.70	2266.	0	2	75166

`tbl_regression()`

```
model_02 |>
  tbl_regression(intercept = TRUE)
```

```
model_02 |>
  tbl_regression(
    intercept = TRUE) |>
  modify_column_hide(conf.low) |>
  modify_column_unhide(column = std.error)
```

OLS Regression predicting number of children by survey year and years of education

Characteristic	Beta	SE	p-value
(Intercept)	5.3	0.813	<0.001
GSS year for this respondent	0.00	0.000	0.045
educ	-0.13	0.002	<0.001

Abbreviations: CI = Confidence Interval, SE = Standard Error

Characteristic	Beta [†]	SE
(Intercept)	5.3***	0.813
GSS year for this respondent	0.00*	0.000
educ	-0.13***	0.002
R ²	0.057	
No. Obs.	75,166	

[†]p<0.05; **p<0.01; ***p<0.001

```
model_02 |>
  tbl_regression(
    intercept = TRUE) |>
  modify_column_hide(conf.low) |>
  modify_column_unhide(column = std.error) |>
  remove_abbreviation() |>
  add_significance_stars() |>
  add_glance_table(include = c("r.squared", "nobs"))
```

1 continuous variable & 1 categorical variable

Are weekly hours worked related to education and gender?

```
model <- lm(hrs1 ~ educ + sex, data = my_data)

get_regression_table(model)
```

```
# A tibble: 3 x 7
  term      estimate std_error statistic p_value lower_ci upper_ci
<chr>      <dbl>    <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 intercept  38.6      0.32    121.     0     38.0    39.2
```


Characteristic	Beta ¹	SE
(Intercept)	39***	0.320
educ	0.42***	0.022
sex		
male	—	—
female	-6.6***	0.132
R ²	0.061	
No. Obs.	43,167	

¹*p<0.05; **p<0.01; ***p<0.001

2 educ	0.42	0.022	18.7	0	0.376	0.464
3 sex: female	-6.58	0.132	-49.9	0	-6.83	-6.32

```
model |>
  tbl_regression(
    intercept = TRUE) |>
  modify_column_hide(conf.low) |>
  modify_column_unhide(
    column = std.error) |>
  remove_abbreviation() |>
  add_significance_stars() |>
  add_glance_table(
    include = c("r.squared", "nobs"))
```

OLS Regression predicting weekly work hours by years of education and gender

```
model |>
  tbl_regression(
    intercept = TRUE,
    show_single_row = sex,
    label = list(sex = "Women")) |>
  modify_column_hide(conf.low) |>
  modify_column_unhide(
    column = std.error) |>
  remove_abbreviation() |>
  add_significance_stars() |>
  add_glance_table(
    include = c("r.squared", "nobs"))
```

Characteristic	Beta ^I	SE
(Intercept)	39***	0.320
educ	0.42***	0.022
Women	-6.6***	0.132
R ²	0.061	
No. Obs.	43,167	

^I*p<0.05; **p<0.01; ***p<0.001

`estimate_means()`

R calculates the mean values of every variable in the model except sex, and sets all of those variables equal to their means. Then (once again) it sets female = 0 in every observation, calculates predicted probabilities for each individual, and averages those.

Then it goes back and resets female to 1 in every observation, again calculates predicted probabilities for each individual and outputs the average of those.

Outputs can be interpreted as the predicted probabilities for women and men if they were all average in all other respects.

```
# Get predicted values
preds <- estimate_means(model, by = c("sex"))

# Preview the new df
preds |>
  as_tibble() |>
  head()
```

```
# A tibble: 2 x 7
  sex      Mean    SE CI_low CI_high    t    df
<fct> <dbl> <dbl> <dbl> <dbl> <dbl> <int>
1 male   44.4 0.0923  44.2   44.5  480. 43164
2 female 37.8 0.0939  37.6   38.0  402. 43164
```

```
# Get predicted values
preds <- estimate_means(model, by = c("educ", "sex"))

# Preview the new df
preds |>
  as_tibble() |>
  head()
```


men with 0 years of education = $38.601 + .420_{educ} * 0$
 women with 0 years of education $38.601 + .420_{educ} * 0 - 6.576_{female}$
 men with 2.222 years of education $38.601 + .420_{educ} * 2.222$
 women with 2.222 years of education $38.601 + .420_{educ} * 2.222 - 6.576_{female}$

Visualize the predictions:

```
# Get predicted values
preds <- estimate_means(model, by = c("educ", "sex"))

# Plot using ggplot2
preds |>
  ggplot(aes(x = educ, y = Mean, color = sex)) +
    geom_line(linewidth = 1.2) +
    labs(
      title = "Predicted Weekly Work Hours by Education and Gender",
      x = "Education Level",
      y = "Predicted Weekly Work Hours",
      color = " " ) +
    theme_minimal() +
    theme(legend.position = "top")
```

Predicted Weekly Work Hours by Education and Gender



Are weekly hours worked related to education and marital status?

```
model <- lm(hrs1 ~ educ + marital, data = my_data)
get_regression_table(model)
```

A tibble: 6 x 7

	term <chr>	estimate <dbl>	std_error <dbl>	statistic <dbl>	p_value <dbl>	lower_ci <dbl>	upper_ci <dbl>
1	intercept	36.5	0.33	110.	0	35.9	37.1
2	educ	0.378	0.023	16.4	0	0.333	0.424
3	marital: widowed	-5.94	0.385	-15.4	0	-6.70	-5.19
4	marital: divorced	0.605	0.198	3.06	0.002	0.218	0.993
5	marital: separated	-0.278	0.379	-0.733	0.464	-1.02	0.465
6	marital: never married	-1.76	0.163	-10.8	0	-2.08	-1.44

Regression Predicting Weekly Work Hours by Years of Education and Marital Status

```
model |>
  tbl_regression(
    intercept = TRUE,
    label = list(
      marital = "Marital status (ref. = married)") |>
  remove_row_type(type = c("reference")) |>
  modify_column_hide(conf.low) |>
  modify_column_unhide(
    column = std.error) |>
  remove_abbreviation() |>
  add_significance_stars() |>
  add_glance_table(
    include = c("r.squared", "nobs"))
```

Are daily TV hours related to education and political identity?

```
my_data <- my_data %>%
  mutate(pol_group = case_when(
    polviews %in% c("extremely liberal", "liberal", "slightly liberal") ~ "Liberal",
    polviews == "moderate, middle of the road" ~ "Moderate",
    polviews %in% c("slightly conservative", "conservative", "extremely conservative") ~ "Conservative",
    TRUE ~ NA_character_))

model <- lm(tvhours ~ educ + pol_group, data = my_data)
```

Characteristic	Beta [†]	SE
(Intercept)	37***	0.330
educ	0.38***	0.023
Marital status (ref. = married)		
widowed	-5.9***	0.385
divorced	0.61**	0.198
separated	-0.28	0.379
never married	-1.8***	0.163
R ²	0.015	
No. Obs.	43,178	

[†]p<0.05; **p<0.01; ***p<0.001

```
get_regression_table(model)
```

```
# A tibble: 4 x 7
```

	term <chr>	estimate <dbl>	std_error <dbl>	statistic <dbl>	p_value <dbl>	lower_ci <dbl>	upper_ci <dbl>
1	intercept	5.00	0.056	89.1	0	4.89	5.12
2	educ	-0.159	0.004	-40.4	0	-0.166	-0.151
3	pol_group: Liberal	0.157	0.031	5.15	0	0.097	0.217
4	pol_group: Moderate	0.192	0.028	6.76	0	0.136	0.247

```
relevel()
```

```
# Relevel pol_group so "Moderate" is the reference
my_data$pol_group <- relevel(factor(my_data$pol_group),
                                ref = "Moderate")

model <- lm(tvhours ~ educ + pol_group, data = my_data)
get_regression_table(model)
```

```
# A tibble: 4 x 7
```

	term <chr>	estimate <dbl>	std_error <dbl>	statistic <dbl>	p_value <dbl>	lower_ci <dbl>	upper_ci <dbl>
1	intercept	5.20	0.054	96.4	0	5.09	5.30
2	educ	-0.159	0.004	-40.4	0	-0.166	-0.151

3	pol_group: Conservative	-0.192	0.028	-6.76	0	-0.247	-0.136
4	pol_group: Liberal	-0.035	0.03	-1.15	0.25	-0.093	0.024

Does the relationship between TV hours and education differ by political identity?

```
model <- lm(tvhours ~ educ * pol_group, data = my_data)

get_regression_table(model)
```

```
# A tibble: 6 x 7
  term                estimate std_error statistic p_value lower_ci upper_ci
  <chr>                <dbl>    <dbl>    <dbl>   <dbl>   <dbl>   <dbl>
1 intercept            4.98      0.088     56.5     0       4.81    5.16
2 educ              -0.142      0.007    -21.1     0      -0.155 -0.129
3 pol_group: Conservative -0.1      0.128    -0.781   0.435   -0.351  0.151
4 pol_group: Liberal    0.553     0.132     4.20     0       0.295  0.811
5 educ:pol_groupConserva~ -0.008     0.01    -0.787   0.431   -0.026  0.011
6 educ:pol_groupLiberal  -0.044     0.01    -4.55     0      -0.063 -0.025
```

The coefficient for `educ` is -0.142 , meaning that for moderates (the reference group), each additional year of education is associated with -0.142 fewer tv hours per day.

The coefficient for `pol_group: Conservative` is -0.1 . This is the difference in the outcome between moderates and conservatives when education = 0. Conservatives watch -0.1 TV hours per day fewer than moderates when education is zero. Not statistically significantly different.

The interaction term `educ:pol_groupConservative` is -0.008 . It tells us whether the relationship between education and TV hours differs between moderates and conservatives. It adjusts the slope of education for conservatives. So for conservatives, the effect of education is:

$$-0.142 - 0.008 = -0.15$$

This means that for conservatives, each additional unit of education is associated with only -0.15 fewer daily TV hours.

Visualize predictions:

```
# Get predicted values
preds <- estimate_means(model, by = c("educ", "pol_group"))

# Plot using ggplot2
preds |>
```

```
ggplot(aes(x = educ, y = Mean, color = pol_group)) +
  geom_line(linewidth = 1.2) +
  labs(
    title = "Predicted Daily TV Hours by Education and Political Identity",
    x = "Education Level",
    y = "Predicted TV Hours",
    color = "Political Identity") +
  theme_minimal() +
  theme(legend.position = "top")
```

Predicted Daily TV Hours by Education and Political Identity

